


Bayesian Analysis of Covid Data

Siddhesh Bagwe

2023-03-09

```
ireland <- read.csv("./ireland1.txt")
cases <- ireland[1:100,1]
days = 1:100
plot(days, cases, xlab="Days in", ylab="Number of daily cases", main="Irish daily covid case numbers")
```

Irish daily covid case numbers



The plot shows the cases over the days. We see that the number of cases kept on increasing with the days. After about the 50-60 day the number of cases started dropping. We will a Bayesian model for the same to predict the number of cases.

```

set.seed(123)
library(rstan)
write("
data {
  int<lower=1> N;
  int y[N];
  vector[N] t;
}

parameters {
  real<lower=0> theta1;
  real<lower=0,upper=100> theta2;
  real<lower=0,upper=1> theta3;
}

model {
  target += normal_lpdf(theta1 | 1e3,1e5);
  target += normal_lpdf(theta2| 50, 100);
  for (n in 1:N){
    target += poisson_lpmf(y[n] | theta1 * theta3 * exp(-theta3 * (t[n] - theta2)) / pow(1 +
exp(-theta3 * (t[n] - theta2)), 2));
  }
}

generated quantities {
  vector[N] y_pred;
  real log_lik[N];
  for (n in 1:N){
    y_pred[n] = poisson_rng( theta1 * theta3 * exp(-theta3 * (t[n] - theta2)) / pow(1 + exp(-
theta3 * (t[n] - theta2)), 2.0) );
    log_lik[n]=poisson_lpmf(y[n] | theta1 * theta3 * exp(-theta3 * (t[n] - theta2)) / pow(1 +
exp(-theta3 * (t[n] - theta2)), 2));
  }
}"
,"m1.stan")

```

Model Description

The above model is defined for the logistic function

$$y_t \sim \text{Po}(\text{lambda}_1(t))$$

- θ_1 , θ_2 are normally distributed while θ_3 is uniformly distributed.
- We will use the y_pred generated to get the prediction from the model. The same can be used to get the accuracy of the model.
- The \log_lik function is defined to compare the accuracy of the model.

```
set.seed(123)
data <- list(N = NROW(cases),
            y = cases,
            t = days)

fit <- stan(file="m1.stan", data = data, iter=500)
```

- The model is fit for 500 iterations using the given data.
- Here t is the number of days and y is number of cases.

```
print(fit, pars=c("theta1", "theta2", "theta3", "y_pred[1]"))
```

```
## Inference for Stan model: anon_model.
## 4 chains, each with iter=500; warmup=250; thin=1;
## post-warmup draws per chain=250, total post-warmup draws=1000.
##
##               mean se_mean      sd      2.5%      25%      50%      75%      97.5%
## theta1      25525.85     5.49 152.67 25239.09 25429.13 25522.25 25619.72 25830.41
## theta2        45.94     0.00   0.11   45.72   45.86   45.94   46.02   46.17
## theta3         0.10     0.00   0.00    0.10    0.10    0.10    0.10    0.10
## y_pred[1]    29.07     0.17   5.43   19.00   25.00   29.00   32.25   40.00
##               n_eff Rhat
## theta1         774 1.01
## theta2        1284 1.00
## theta3         827 1.00
## y_pred[1]      971 1.00
##
## Samples were drawn using NUTS(diag_e) at Mon May 29 22:02:57 2023.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```

Model Summary

Here we see the summary of our model. - The model is fit with 4 chains each with 500 iterations. Out of this 250 are used as warm-up. Thus we get 250 post-warm up draws per chain giving a total of 1000 post-warm up draws. - The parameters theta1, theta2 and theta3 define the function in our model. We see the mean, Standard error, standard deviation as well as the quantiles for our parameters. These parameters are then used to get the prediction from the model. As the data is of 100 days, we have the y_pred from 1 to 100. Here we see the mean, standard deviation as well as the quantiles for prediction of day 1(y_pred[1]). - n_eff is the effective sample_size. A sample size of greater than 10% of the total post-warm draws is acceptable (100 for this model.). The summary shows that all the parameters have an acceptable n_eff. - R_hat is the measure of convergence of the chains. The acceptable value of Rhat is less than or equal to 1.1. As all the parameters have Rhat less than 1.1 the model is acceptable.

```
y_pred <- as.matrix(fit, pars=c("y_pred"))
y_phdi = HDInterval::hdi(y_pred, credMass=0.90)
pi_l = y_phdi[1,]
pi_u = y_phdi[2,]
print(y_phdi)
```

```

##      parameters
##      y_pred[1] y_pred[2] y_pred[3] y_pred[4] y_pred[5] y_pred[6] y_pred[7]
## lower      21      22      24      29      31      34      37
## upper      38      41      43      48      52      57      61
##      parameters
##      y_pred[8] y_pred[9] y_pred[10] y_pred[11] y_pred[12] y_pred[13]
## lower      42      49      53      60      67      76
## upper      67      75      81      88      96     106
##      parameters
##      y_pred[14] y_pred[15] y_pred[16] y_pred[17] y_pred[18] y_pred[19]
## lower      81      90      99     108     119     132
## upper     113     125     135     146     158     173
##      parameters
##      y_pred[20] y_pred[21] y_pred[22] y_pred[23] y_pred[24] y_pred[25]
## lower     148     158     176     189     207     224
## upper     189     202     223     235     258     275
##      parameters
##      y_pred[26] y_pred[27] y_pred[28] y_pred[29] y_pred[30] y_pred[31]
## lower     242     262     280     304     325     347
## upper     297     317     341     365     388     412
##      parameters
##      y_pred[32] y_pred[33] y_pred[34] y_pred[35] y_pred[36] y_pred[37]
## lower     373     395     422     438     463     485
## upper     440     462     493     511     537     562
##      parameters
##      y_pred[38] y_pred[39] y_pred[40] y_pred[41] y_pred[42] y_pred[43]
## lower     505     520     539     550     570     575
## upper     578     599     618     630     648     654
##      parameters
##      y_pred[44] y_pred[45] y_pred[46] y_pred[47] y_pred[48] y_pred[49]
## lower     582     588     584     587     583     581
## upper     667     670     668     671     663     659
##      parameters
##      y_pred[50] y_pred[51] y_pred[52] y_pred[53] y_pred[54] y_pred[55]
## lower     566     554     536     523     502     484
## upper     647     631     614     599     576     556
##      parameters
##      y_pred[56] y_pred[57] y_pred[58] y_pred[59] y_pred[60] y_pred[61]
## lower     462     436     418     392     370     349
## upper     534     506     490     459     434     412
##      parameters
##      y_pred[62] y_pred[63] y_pred[64] y_pred[65] y_pred[66] y_pred[67]
## lower     319     303     279     261     239     220
## upper     381     363     338     317     296     273
##      parameters
##      y_pred[68] y_pred[69] y_pred[70] y_pred[71] y_pred[72] y_pred[73]
## lower     207     188     173     155     143     132
## upper     256     235     220     199     186     171
##      parameters
##      y_pred[74] y_pred[75] y_pred[76] y_pred[77] y_pred[78] y_pred[79]
## lower     122     108     100      88      82      72
## upper     159     145     136     122     113     104
##      parameters
##      y_pred[80] y_pred[81] y_pred[82] y_pred[83] y_pred[84] y_pred[85]
## lower      66      59      52      47      43      37

```

```
##      upper      96      88      79      74      67      61
##      parameters
##      y_pred[86] y_pred[87] y_pred[88] y_pred[89] y_pred[90] y_pred[91]
##      lower      35      31      28      25      21      20
##      upper      57      52      46      44      40      37
##      parameters
##      y_pred[92] y_pred[93] y_pred[94] y_pred[95] y_pred[96] y_pred[97]
##      lower      16      16      13      11      11      8
##      upper      33      31      28      25      24      21
##      parameters
##      y_pred[98] y_pred[99] y_pred[100]
##      lower      8      8      6
##      upper      20      19      17
## attr(,"credMass")
## [1] 0.9
```


We can see the 90% posterior interval of our model. We will use this to check the accuracy of our model.

```
d1 <- as.data.frame(data)
library(ggplot2)
p <- ggplot()

p2 <- p +
  geom_point(data = d1,
    aes(t, y), shape = 1, color = 'dodgerblue') +
  ggtitle("Prediction Interval = 0.90")+
  geom_ribbon(data = d1,
    mapping = aes(t,ymin=pi_l, ymax=pi_u), alpha = .05,color = 'red')+xlab("Days")+ ylab("C
ases")


p2
```

Prediction Interval = 0.90



The plot shows the ribbon of the 90% interval of our predicted data and the plot of the original data. We can see that though our model does a decent job in predicting the curve, the real values do not lie in the 90% interval for most of our data. This indicates that the model may not be a perfect fit for our data. We will check this further using the density plot of our original data to that of the posterior prediction.

```
library(bayesplot)
ppc_dens_overlay(d1$y, y_pred[1:100,])+ theme_classic()
```



Again we can see that, though the curve is similar there is quite a difference between the original and predicted values. This indicates that the model may not be a great fit.

Next we will try to use our model to make the predication for number of cases for the next five days.

```

set.seed(123)
library(rstan)
write("
data {
  int<lower=1> N;
  int<lower=1> N1;
  int y[N];
  vector[N] t;
  vector[N1] t_new;
}

parameters {
  real<lower=0> theta1;
  real<lower=0,upper=100> theta2;
  real<lower=0,upper=1> theta3;
}

model {
  target += normal_lpdf(theta1 | 1e3,1e5);
  target += normal_lpdf(theta2| 50, 100);
  for (n in 1:N){
    target += poisson_lpmf(y[n] | theta1 * theta3 * exp(-theta3 * (t[n] - theta2)) / pow(1 +
exp(-theta3 * (t[n] - theta2)), 2));
  }
}

generated quantities {
  vector[N1] y_pred;
  for (n in 1:N1){
    y_pred[n] = poisson_rng( theta1 * theta3 * exp(-theta3 * (t_new[n] - theta2)) / pow(1 + e
xp(-theta3 * (t_new[n] - theta2)), 2.0) );
  }
}"
,"m2.stan")

```

Here the model is updated to predict values for `t_new` which is for days 101-105. We have to define the new data to fit our model.

```

set.seed(123)
t_new <- c(101,102,103,104,105)
data1 <- list(N = NROW(cases),
             N1 = as.integer(5),
             y = cases,
             t = days,
             t_new=t_new)

fit1 <- stan(file="m2.stan", data = data1, iter=500)

```

- The model is fit again with new data
- Here along with the days and cases, we also give the data `t_new` to get the prediction.


```
y_pred1 <- extract(fit1)$y_pred
cases_pred <- as.integer(c(mean(y_pred1[,1]),mean(y_pred1[,2]),mean(y_pred1[,3]),mean(y_pred1[,4]),mean(y_pred1[,5])))
print(cases_pred)
```

```
## [1] 10 9 8 8 7
```

Our model has predicted the number of cases for the next five days would be 10, 9, 8, 8, 7 respectively. Although the prediction may not be accurate, it can still give us a rough idea of the number of cases that can be found over the next five days.

Again we will create our model to predict the number of cases but this time we will use the $g(t)$ function. The $g(t)$ function is given as

$$g(t) = \theta_1 \exp(-\theta_2 \theta_3^t)$$

For our model we need the value of $\lambda_g(t)$ which is the derivative of $g(t)$ w.r.t t .

Taking derivative we get $\lambda_g(t)$ as

$$\lambda_g(t) = \theta_1 * -\theta_2 * \theta_3^t * \exp(-\theta_2 * \theta_3^t) * \log(\theta_3)$$

We will use this function to build our model

```

set.seed(123)
library(rstan)
write("
data {
  int<lower=1> N;
  int y[N];
  vector[N] t;
}

parameters {
  real<lower=0> theta1;
  real<lower=0,upper=100> theta2;
  real<lower=0,upper=1> theta3;
}

model {
  target += normal_lpdf(theta1 | 1e3,1e5);
  target += normal_lpdf(theta2| 50, 100);
  for (n in 1:N){
    target += poisson_lpmf(y[n] | theta1 * (-theta2) * pow(theta3,t[n]) * exp(-theta2 * pow(theta3,t[n]))) * log(theta3));
  }
}

generated quantities {
  vector[N] y_pred;
  real log_lik[N];
  for (n in 1:N){
    y_pred[n] = poisson_rng( theta1 * (-theta2) * pow(theta3,t[n]) * exp(-theta2 * pow(theta3,t[n])) * log(theta3));
    log_lik[n] = poisson_lpmf(y[n] | theta1 * (-theta2) * pow(theta3,t[n]) * exp(-theta2 * pow(theta3,t[n])) * log(theta3));
  }
}"
,"m3.stan")

fit2<- stan(file="m3.stan", data = data, iter=500)

```

Our model is trained for 500 iterations using the same data as for part 1.

```
print(fit2, pars=c("theta1", "theta2", "theta3","y_pred[1]"))
```

```
## Inference for Stan model: anon_model.
## 4 chains, each with iter=500; warmup=250; thin=1;
## post-warmup draws per chain=250, total post-warmup draws=1000.
##
##               mean se_mean      sd      2.5%      25%      50%      75%      97.5%
## theta1      25558.68     6.45 161.38 25252.69 25448.54 25556.77 25666.33 25869.51
## theta2       14.19     0.01   0.19   13.80   14.08   14.20   14.32   14.54
## theta3        0.94     0.00   0.00    0.93    0.94    0.94    0.94    0.94
## y_pred[1]    0.03     0.01   0.17    0.00    0.00    0.00    0.00    1.00
##               n_eff Rhat
## theta1        626 1.00
## theta2        375 1.01
## theta3        387 1.00
## y_pred[1]     797 1.00
##
## Samples were drawn using NUTS(diag_e) at Mon May 29 22:04:27 2023.
## For each parameter, n_eff is a crude measure of effective sample size,
## and Rhat is the potential scale reduction factor on split chains (at
## convergence, Rhat=1).
```

-The model summary can be seen using the print function. We get the mean, sd and the quantiles for the parameters theta1, theta2, theta3 as well as the predictions. We have shown the data for the prediction of first day. - The noticeable difference is the summary of y_pred[1] of both the models. In this model we get that the lower and upper levels for y_pred[1] is 0 to 1 with a mean of 0.03 and sd of 0.17. - The n_eff and Rhat values are acceptable for this model as well.

```
y_pred2 <- as.matrix(fit2, pars=c("y_pred"))
y_phdi1 = HDInterval::hdi(y_pred2, credMass=0.90)
pi_l1 = y_phdi1[1,]
pi_u1 = y_phdi1[2,]
print(y_phdi1)
```

```

##      parameters
##      y_pred[1] y_pred[2] y_pred[3] y_pred[4] y_pred[5] y_pred[6] y_pred[7]
## lower      0      0      0      0      0      0      0
## upper      0      0      1      1      2      3      4
##      parameters
##      y_pred[8] y_pred[9] y_pred[10] y_pred[11] y_pred[12] y_pred[13]
## lower      1      2      4      7      11      18
## upper      6      9      13      18      25      34
##      parameters
##      y_pred[14] y_pred[15] y_pred[16] y_pred[17] y_pred[18] y_pred[19]
## lower      27      35      50      68      87      107
## upper      47      59      76      98      120      144
##      parameters
##      y_pred[20] y_pred[21] y_pred[22] y_pred[23] y_pred[24] y_pred[25]
## lower      131      161      189      224      254      287
## upper      171      206      237      276      310      348
##      parameters
##      y_pred[26] y_pred[27] y_pred[28] y_pred[29] y_pred[30] y_pred[31]
## lower      318      355      386      412      445      477
## upper      381      420      456      484      519      555
##      parameters
##      y_pred[32] y_pred[33] y_pred[34] y_pred[35] y_pred[36] y_pred[37]
## lower      500      523      541      555      566      579
## upper      576      600      619      634      648      657
##      parameters
##      y_pred[38] y_pred[39] y_pred[40] y_pred[41] y_pred[42] y_pred[43]
## lower      580      585      587      581      578      574
## upper      662      669      668      667      663      656
##      parameters
##      y_pred[44] y_pred[45] y_pred[46] y_pred[47] y_pred[48] y_pred[49]
## lower      566      555      540      527      516      496
## upper      644      633      617      605      593      571
##      parameters
##      y_pred[50] y_pred[51] y_pred[52] y_pred[53] y_pred[54] y_pred[55]
## lower      480      468      447      426      414      392
## upper      555      539      522      496      484      459
##      parameters
##      y_pred[56] y_pred[57] y_pred[58] y_pred[59] y_pred[60] y_pred[61]
## lower      375      356      337      325      304      291
## upper      439      426      402      387      364      348
##      parameters
##      y_pred[62] y_pred[63] y_pred[64] y_pred[65] y_pred[66] y_pred[67]
## lower      277      260      248      235      218      208
## upper      334      316      303      287      270      256
##      parameters
##      y_pred[68] y_pred[69] y_pred[70] y_pred[71] y_pred[72] y_pred[73]
## lower      196      182      175      163      151      143
## upper      243      230      221      205      197      186
##      parameters
##      y_pred[74] y_pred[75] y_pred[76] y_pred[77] y_pred[78] y_pred[79]
## lower      134      125      121      113      102      96
## upper      175      165      160      150      139      129
##      parameters
##      y_pred[80] y_pred[81] y_pred[82] y_pred[83] y_pred[84] y_pred[85]
## lower      89      85      78      70      70      64

```


```
##      upper      123      118      111      102      99      93
##      parameters
##      y_pred[86] y_pred[87] y_pred[88] y_pred[89] y_pred[90] y_pred[91]
##      lower      57      55      52      46      43      41
##      upper      86      82      78      71      68      64
##      parameters
##      y_pred[92] y_pred[93] y_pred[94] y_pred[95] y_pred[96] y_pred[97]
##      lower      38      35      33      30      28      26
##      upper      60      57      54      50      48      45
##      parameters
##      y_pred[98] y_pred[99] y_pred[100]
##      lower      24      22      20
##      upper      42      40      38
## attr(,"credMass")
## [1] 0.9
```

The 90% credible intervals shows that the model predicts the cases to fall between this range. We can clearly see the difference between this model and the model in part1. Plotting the credible intervals with the original data will give us an idea of the accuracy of our model.

```
p3 <- p +
  geom_point(data = d1,
    aes(t, y), shape = 1, color = 'dodgerblue') +
  ggtitle("Prediction Interval = 0.90")+
  geom_ribbon(data = d1,
    mapping = aes(t, ymin=pi_l1, ymax=pi_u1), alpha = .05,color = 'red')+xlab("Days")+ ylab
("Cases")

p3
```


Prediction Interval = 0.90



From the plot we can again see that the model doesn't do a great job in predicting the values. Although the values are predicted accurately at the start and to the end but there is a lot of inconsistency in the centre.

```
library(bayesplot)
ppc_dens_overlay(d1$y, y_pred2[1:100,])+ theme_classic()
```

```
## Warning: Ignoring unknown parameters: linewidth
## Ignoring unknown parameters: linewidth
```



Through bayesplot, we can see the density curve of the original data as compared to our prediction.

```

set.seed(123)
library(rstan)
write("
data {
  int<lower=1> N;
  int<lower=1> N1;
  int y[N];
  vector[N] t;
  vector[N1] t_new;
}

parameters {
  real<lower=0> theta1;
  real<lower=0,upper=100> theta2;
  real<lower=0,upper=1> theta3;
}

model {
  target += normal_lpdf(theta1 | 1e3,1e5);
  target += normal_lpdf(theta2| 50, 100);
  for (n in 1:N){
    target += poisson_lpmf(y[n] | theta1 * (-theta2) * pow(theta3,t[n]) * exp(-theta2 * pow(theta3,t[n])) * log(theta3));
  }
}

generated quantities {
  vector[N1] y_pred;
  for (n in 1:N1){
    y_pred[n] = poisson_rng( theta1 * (-theta2) * pow(theta3,t_new[n]) * exp(-theta2 * pow(theta3,t_new[n])) * log(theta3) );
  }
}"
,"m4.stan")

fit3 <- stan(file="m4.stan", data = data1, iter=500)

```

```

y_pred3 <- extract(fit3)$y_pred
cases_pred1 <- as.integer(c(mean(y_pred3[,1]),mean(y_pred3[,2]),mean(y_pred3[,3]),mean(y_pred3[,4]),mean(y_pred3[,5])))
print(cases_pred1)

```

```
## [1] 27 25 24 22 21
```

We see that the prediction from the model for the days 101-105 is 27, 25, 24, 22, 21 respectively. As the model gives 500 iterations, we use the mean of 500 values to get the prediction.

Model Comparison using Loo package

```
library(loo)
```



```
## This is loo version 2.5.1
```

```
## - Online documentation and vignettes at mc-stan.org/loo
```

```
## - As of v2.0.0 loo defaults to 1 core but we recommend using as many as possible. Use the
'cores' argument or set options(mc.cores = NUM_CORES) for an entire session.
```

```
## - Windows 10 users: loo may be very slow if 'mc.cores' is set in your .Rprofile file (see
https://github.com/stan-dev/loo/issues/94).
```

```
##
## Attaching package: 'loo'
```

```
## The following object is masked from 'package:rstan':
##
##      loo
```

```
loo1 <- loo(fit)
loo2 <- loo(fit2)
loo_compare(loo1, loo2)
```

```
##          elpd_diff se_diff
## model2      0.0      0.0
## model1 -922.3    248.0
```

From the values of `loo_compare` we see that the accuracy of model 2(Gompertz) is better than that of the logistic function. So model 2 is preferred over model 1. The standard error difference too is pretty high, indicating that there is a big difference between the two. We will use the WAIC function to get more information.

Model Comparison using WAIC function

```
log_lik1 <- loo::extract_log_lik(fit)
waic1 <- loo::waic(log_lik1)
log_lik2 <- loo::extract_log_lik(fit2)
waic2 <- loo::waic(log_lik2)
loo_compare(waic(log_lik1), waic(log_lik2))
```

```
##          elpd_diff se_diff
## model2      0.0      0.0
## model1 -926.3    249.4
```

The WAIC function also suggests that model 2 is better than model 1. The standard in this case is also high.

Although model 2 is better than model 1, both the models are notable to predict the data as accurately as desired. We can improve the model fit by the following techniques

Methods to improve model fit

1. We can adjust the model parameters (θ_1 , θ_2 , θ_3) to improve the fit.
2. We can add more data to improve the model. More data helps the model to learn the trend better so as to better estimate the values.
3. We can try different models and check whether a different model can be a better fit.
4. Regularization techniques can help to prevent overfitting of the data by adding a penalty term to the model, which can improve the fit.

Conclusion

Overall we see that the Gompertz model is a better model for our data but it can be eventually be improved.